

Networking Architecture — LAN Share

This document explains the project from a **networking perspective**, mapping every piece of code to the textbook concept it implements. Read it alongside [protocol.py](#), [discovery.py](#), and [transfer.py](#).

1. The OSI / TCP-IP layers used

Application layer	JSON-framed protocol (our code)
Transport layer	TCP (file transfer) UDP (device discovery)
Network layer	IP
Link layer	Ethernet / Wi-Fi (handled by OS)

We write code at the **Application layer**. The OS handles everything below.

2. Why two transport protocols?

Concern	Discovery	File transfer
Recipient known in advance?	No	Yes (chosen from peer list)
Reliability needed?	No (next HELLO replaces lost ones)	Yes — every byte must arrive
Ordering needed?	No	Yes — chunk N before chunk N+1
Connection setup OK?	No (connection-less, no handshake)	Yes (3-way handshake is fine once)
Choice	UDP	TCP

This is the classic textbook split: **UDP for announcements, TCP for streams**.

3. Discovery: UDP multicast

Why multicast (239.42.42.42), not broadcast (255.255.255.255)?

UDP broadcast is the textbook example, and it works on most LANs. We use **multicast** for one practical reason: **on Windows, when two processes on the same machine both bind UDP port 5002, only one receives broadcasted packets** — broadcast delivery is undefined when multiple sockets share a port. Multicast delivery, in contrast, is **duplicated to every socket that has joined the group**, so two processes on the same PC reliably discover each other.

Multicast group **239.42.42.42** is in the **locally-scoped administrative range** (**239.0.0.0/8**, see [RFC 2365](#)). With **TTL=1** the packets travel exactly one hop — never crossing routers, staying on the local LAN.

The three threads inside **Discovery** ([discovery.py](#))

```
flowchart LR
    Beacon["beacon_loop\n(every 3s)"] -->| "sendto(239.42.42.42:5002)" | Net["Net((LAN))"]
    Net -->| "recvfrom" | Listener["listener_loop\n(blocking)"]
    Listener --> PeerDict["self.peers\n(device_id → Peer)"]
    Sweeper["sweeper_loop\n(every 2s)"] -->| "drop if last_seen > 10s ago" | PeerDict
    PeerDict
```

The **HELLO** packet (one UDP datagram)

```
{
  "type":      "HELLO",
  "device_id": "5f3c7b...", // stable per-machine UUID
  "name":      "Ahmad-Laptop", // friendly name shown in UI
  "tcp_port":  5001,          // where to TCP-connect to send files
  "version":   2
}
```

Key socket calls (with line references)

Sender side ([discovery.py _make_sender](#)):

```
s = socket.socket(AF_INET, SOCK_DGRAM)           # UDP socket
s.setsockopt(IPPROTO_IP, IP_MULTICAST_TTL, 1)    # local LAN only
s.setsockopt(IPPROTO_IP, IP_MULTICAST_LOOP, 1)   # see our own packets
s.sendto(packet, ("239.42.42.42", 5002))         # ← send the datagram
```

Receiver side ([discovery.py _make_listener](#)):

```
s = socket.socket(AF_INET, SOCK_DGRAM)
s.setsockopt(SOL_SOCKET, SO_REUSEADDR, 1)        # share port with siblings
s.bind("", 5002)                                  # listen on all interfaces
mreq = struct.pack("4s1", inet_aton("239.42.42.42"), INADDR_ANY)
s.setsockopt(IPPROTO_IP, IP_ADD_MEMBERSHIP, mreq) # ← join multicast group
data, addr = s.recvfrom(4096)                    # ← receive datagram
```

IP_ADD_MEMBERSHIP tells the kernel: "I want to receive packets sent to multicast group X on any interface." Without it, the kernel filters multicast packets out.

Peer expiry (TTL)

Peers have a `last_seen` timestamp. The `sweeper_loop` runs every 2 seconds and removes entries older than `PEER_TTL=10` seconds. So when a laptop closes its lid, everyone notices within 10 seconds — even though no `BYE` was sent.

4. File transfer: TCP

Why we still need application-layer logic on top of TCP

TCP gives us:

- **Reliability** — retransmission, error correction at packet level
- **Ordering** — bytes arrive in order
- **Flow control** — sender slows down when receiver is full
- **Congestion control** — backs off when network is busy

TCP does **not** give us:

- Message boundaries — TCP is a **byte stream**, not a message stream
- Application-level structure — what does each byte mean?
- End-to-end integrity — TCP's checksum is only 16 bits per segment, can be defeated by a hash collision or corruption after the kernel hands data over

So our application layer adds:

1. **Message framing** (`protocol.py`) — 4-byte length + JSON header + optional binary payload
2. **Typed messages** — `REQUEST`, `RESPONSE`, `METADATA`, `DATA`, `ACK`, `DONE`, ...
3. **A handshake** — `REQUEST/RESPONSE` so the user can accept or reject
4. **Per-chunk ACKs** — fine-grained progress + early failure detection
5. **End-to-end SHA-256** — catches every kind of corruption, including disk errors

The full session timeline

```
sequenceDiagram
    participant S as Sender
    participant R as Receiver

    Note over S,R: TCP three-way handshake
    S->>R: SYN
    R->>S: SYN+ACK
    S->>R: ACK

    Note over S,R: Application protocol begins
    S->>R: REQUEST { manifest:[...], total_size }
    Note right of R: UI shows accept dialog
    R->>S: RESPONSE { accepted: true }

    loop for each file
        S->>R: METADATA { filename, size, sha256 }
        R->>S: ACK ok
        loop for each 256 KB chunk
```

```

        S->>R: DATA { payload_size + raw bytes }
        R->>S: ACK ok
    end
    S->>R: DONE
    Note right of R: Verify SHA-256
    R->>S: ACK ok
end

S->>R: SESSION_DONE
R->>S: ACK ok

Note over S,R: TCP four-way close
S->>R: FIN
R->>S: FIN+ACK

```

Key socket calls

Sender ([transfer.py](#) [send_files](#)):

```

sock = socket.socket(AF_INET, SOCK_STREAM)           # TCP socket
sock.connect((host, port))                          # ← TCP three-way handshake
send_msg(sock, REQUEST, manifest=..., ...)          # send JSON-framed message
resp = recv_msg(sock)                               # block until RESPONSE arrives
for chunk in iter_chunks(file):
    send_msg(sock, DATA, payload=chunk, ...)        # send one DATA message
    ack = recv_msg(sock)                            # wait for ACK
sock.close()                                         # ← FIN, TCP four-way close

```

Receiver ([transfer.py](#) [serve_forever](#)):

```

server = socket.socket(AF_INET, SOCK_STREAM)
server.bind("", 5001)                                # INADDR_ANY: all interfaces
server.listen(8)                                     # backlog: queued half-open
conns
while True:
    conn, addr = server.accept()                     # ← block until client
    connects
    threading.Thread(target=handle, args=(conn,)).start()

```

`bind("", 5001)` uses `INADDR_ANY` (`0.0.0.0`) so we accept connections on **any** local interface (Wi-Fi, Ethernet, loopback). `listen(8)` sets the kernel's accept backlog — up to 8 pending connections can wait while we're busy.

Why per-chunk ACK?

Without per-chunk ACK, the sender would happily push 100 MB into the kernel's TCP send buffer, then think it's "done" — even if the receiver crashed after byte 1. Per-chunk ACK is round-trip-bound, so the sender

knows after each 256 KB whether the receiver is still alive and writing to disk successfully. It also gives smooth UI progress (1 chunk = 1 progress update).

The trade-off: **throughput drops**. A single round-trip-time (RTT) per chunk caps speed. On a 1ms-RTT LAN with 256 KB chunks, the upper bound is ~250 MB/s — fine for our use case. On a 50ms-RTT WAN it would be ~5 MB/s, which is why production protocols (HTTP/2, gRPC, BitTorrent) **pipeline** multiple chunks before waiting for ACKs.

Why end-to-end SHA-256?

Because **TCP's 16-bit checksum can be defeated**. Some real failure modes:

- A faulty NIC corrupts bytes after TCP checksum verification
- The kernel writes correctly to disk, but the disk has a bad sector
- Memory corruption between socket buffer and `f.write()`

SHA-256 is computed by the **sender** before transmission and verified by the **receiver** after writing to disk. If a single bit differs anywhere along the entire path, the digest changes and we delete the corrupted file. This is the "end-to-end argument in system design" (Saltzer, Reed, Clark — 1984).

5. Concurrency model

The app uses **OS threads**, not async I/O, because socket programming with threads is the textbook model and easier to reason about for a course discussion.

```

flowchart TB
    subgraph MainProcess [LAN Share Process]
        UIThread["Main thread  
(Tkinter UI)"]
        BeaconT["beacon thread  
send HELLO every 3s"]
        ListenerT["listener thread  
UDP recvfrom"]
        SweeperT["sweeper thread  
expire stale peers"]
        AcceptT["TCP accept thread  
server.accept()"]
        RecvT1["receiver handler 1  
(spawned per connection)"]
        RecvT2["receiver handler 2  
(spawned per connection)"]
        SendT["sender thread  
(spawned per outgoing transfer)"]
    end

    UIThread -.->| "queue.Queue" | BeaconT
    UIThread -.->| "queue.Queue" | ListenerT
    UIThread -.->| "queue.Queue" | AcceptT
    AcceptT --> RecvT1
    AcceptT --> RecvT2
  
```

Thread safety

Tkinter is **not** thread-safe — only the main thread may touch widgets. Every other thread sends events to the UI through `queue.Queue`. The UI calls `self.after(50, self._poll_queue)` to drain the queue from the main thread on a 50 ms timer.

6. Reliability summary

Failure mode	What catches it
Single bit flip in a TCP segment	TCP checksum (16-bit)
Lost TCP segment	TCP retransmission
Out-of-order TCP segments	TCP reassembly buffer
Connection drop mid-transfer	<code>recv()</code> returns 0 → <code>ConnectionError</code> raised → UI shows error
Receiver disk full / file write fails	Exception in <code>_handle_connection</code> → ERROR sent → sender aborts
Bit corruption <i>outside</i> the TCP stack (NIC, RAM, disk)	Application-level SHA-256 at end of each file
Receiver rejects the transfer	<code>RESPONSE { accepted: false }</code>
Sender or receiver crashes	TCP RST → other side gets <code>ConnectionResetError</code>
Discovery packet lost	Next HELLO replaces it (every 3 s)
Peer goes offline	Peer entry expires after 10 s of silence

7. What you can demonstrate live

Question the prof might ask	Demo / answer
"Show me the TCP three-way handshake"	Wireshark filter: <code>tcp.port == 5001</code> while clicking Send. Three packets: SYN, SYN+ACK, ACK
"Show me the application-layer messages"	Wireshark "Follow TCP Stream" — you can see the JSON headers in cleartext
"What if you remove SHA-256?"	Comment out the verification block → flip a bit in a received file → silently corrupted, sender thinks it succeeded
"What if you use UDP for transfer?"	Would need to implement reliability (sequence numbers, retransmission, ordering) — i.e. reinvent TCP poorly

Question the prof might ask	Demo / answer
"Why multicast not broadcast?"	Two processes on the same Windows machine sharing UDP port 5002 — multicast delivers to both, broadcast doesn't
"How does discovery survive packet loss?"	Each HELLO is independent; the next one (3 s later) re-announces. TTL=10 s gives 3 chances before a peer expires
"What's the security model?"	Receiver shows a dialog with the file list before any DATA is sent. Currently no encryption — could add TLS via <code>ssl.wrap_socket()</code> in ~5 lines

8. Suggested talking points

1. **"We chose TCP for transfer because we need ordering, reliability, and flow control — these are exactly what TCP provides for free, and reimplementing them on UDP would mean writing a worse TCP."**
2. **"We chose UDP for discovery because it's connectionless. We don't know who the peers are yet, and we don't need reliability — if a HELLO is lost, the next one (3 seconds later) will replace it."**
3. **"TCP gives us a byte stream, not a message stream. Our framing layer (4-byte length prefix + JSON header) recovers message boundaries — without it, `recv()` would return arbitrary fragments."**
4. **"Per-chunk ACK gives us application-layer flow visibility. TCP's flow control happens between kernels; our ACK happens between processes, so we know the receiver actually wrote the chunk to disk."**
5. **"End-to-end integrity needs end-to-end checks. TCP's 16-bit checksum protects only the network path. SHA-256 over the whole file protects against everything — including bugs in our own code."**
6. **"Multicast is a network-layer feature implemented by IGMP. We join group 239.42.42.42 with `IP_ADD_MEMBERSHIP`, the kernel sends an IGMP report, and the local switch starts forwarding multicast frames to our interface. TTL=1 keeps us on the LAN."**